Project Sentinel Project Charter v2.0

Product Philosophy & User Experience Principles

Project: Project Sentinel

Document: Project Charter v2.0 - Product Philosophy & User Experience Principles

Status: Approved for Pre-Coding Baseline

Date: 2026-07-09

Language: English

Scope: Applies to all future Engineering Archives, implementation milestones, user interfaces, APIs, reports, dashboards, automated recommendations, AI explanations, and remediation workflows.

1. Purpose

This Charter update establishes mandatory product and user-experience principles for Project Sentinel.

Project Sentinel shall not only detect security problems; it shall make security understandable, explainable, evidence-based, and actionable for users ranging from home users to enterprise security professionals.

The objective is to ensure that every finding, recommendation, workflow, report, dashboard, and automated action follows a human-centered security model.

2. Mission Statement

Project Sentinel shall provide an intelligent Cyber Security Operating Environment that enables users of all experience levels to understand, evaluate, and improve their security posture through evidence-driven analysis, understandable explanations, and practical remediation guidance.

The platform shall prioritize clarity, trust, transparency, privacy, and usability without sacrificing technical depth or engineering quality.

3. Permanent Product Principle

> **Every finding must be understandable by a non-expert and actionable by an expert.**

This principle is mandatory and applies across the full product lifecycle.

A finding is not considered complete unless it explains what was found, why it matters, how it was determined, which evidence supports it, what should be done next, and what deeper technical details are available.

4. Product Philosophy Principles

Principle 1 - Explain Before You Recommend

Every finding shall clearly explain:

- What happened.
- Why it matters.
- How Sentinel determined the finding.
- Which evidence supports the conclusion.
- What risks exist if no action is taken.
- Which remediation actions are recommended.
- The expected outcome after remediation.

Recommendations shall never appear without an accompanying explanation.

Principle 2 - Simplicity First

The platform shall present security information using language appropriate to the user's experience level.

Complex technical terminology shall only be introduced when necessary and shall always be accompanied by plain-language explanations.

Users shall never be expected to understand cybersecurity terminology in order to understand security risks.

Principle 3 - Evidence Before Opinion

Sentinel shall never generate recommendations based solely on assumptions.

Every recommendation shall reference collected evidence, policy evaluation, engineering rules, threat intelligence, risk calculations, or deterministic analysis.

Evidence shall always be accessible to the user.

Principle 4 - Human-Centered Security

The objective of Project Sentinel is to help people improve security, not to overwhelm them with technical information.

Every interaction should reduce uncertainty and increase confidence.

The platform shall encourage secure behavior through education, understandable guidance, and actionable recommendations.

Principle 5 - Expert Depth on Demand

Every simplified finding shall provide access to progressively deeper technical information.

The user interface shall support multiple information levels without duplicating data.

Progressive Detail Model

- **Level 1 - Summary:** What is the problem?
- **Level 2 - Explanation:** Why does it matter?
- **Level 3 - Evidence:** What data proves it?

- **Level 4 - Technical Detail:** What exact configuration, event, object, or rule caused it?
- **Level 5 - Remediation:** What should be done, manually or automatically?
- **Level 6 - Audit Trail:** What changed and when?

Principle 6 - Transparency by Design

Sentinel shall clearly communicate:

- What data has been collected.
- Why it has been collected.
- Where it is processed.
- How long it is retained.
- Who can access it.
- Whether the analysis is local-only, cloud-assisted, or enterprise-managed.

Users shall never be surprised by platform behavior.

Principle 7 - Privacy First

Privacy is a fundamental design principle.

Project Sentinel shall support:

- Local-only operation.
- Offline analysis where practical.
- User-controlled telemetry.
- Data minimization.
- Explicit consent.
- Encryption by default.
- Clear retention policies.

The platform shall collect only the information required to perform the requested analysis.

Principle 8 - Security Without Fear

Sentinel shall avoid alarmist messaging.

Instead of creating anxiety, the platform shall explain risks objectively, prioritize findings, recommend achievable actions, and measure improvement over time.

Principle 9 - Progressive Guidance

Recommendations shall be prioritized according to:

1. Critical actions.
2. High-value improvements.
3. Best practices.

4. Long-term optimization.

Users should always know what to do next.

Principle 10 - Trust Through Engineering

Every component of Project Sentinel shall be deterministic where possible, evidence-driven, traceable, auditable, reproducible, and explainable.

Trust shall be earned through engineering quality rather than marketing claims.

5. Finding Communication Standard

Every user-facing finding shall use the following standard structure:

Required Finding Fields

Field	Purpose	Required
Title	Plain-language summary	Yes
Severity	Risk level	Yes
What Happened	Direct explanation of the issue	Yes
Why It Matters	Security impact	Yes
Evidence	Supporting data and references	Yes
Affected Assets	Devices, identities, domains, applications, or services	Yes
Recommended Action	Clear next step	Yes
Fix Difficulty	Expected effort	Yes
Automation Eligibility	Whether Sentinel can safely fix it	Yes
Expert Details	Technical expansion	Yes
Audit Trail	History and decisions	Yes

6. Standard User Output Pattern

Sentinel findings shall be rendered using a consistent pattern:

Problem: What Sentinel found.

Risk: Why it matters.

Evidence: What proves it.

Fix: What should be done.

Impact: What changes after the fix.

Details: Expand for technical evidence.

7. Example Finding

Non-Preferred Output

Weak password detected.

Sentinel-Compliant Output

Problem

Your NAS administrator account is using a password that does not meet the security policy.

Why it matters

Administrative accounts can control files, backups, and system settings. If this account is compromised, an attacker may gain broad access to your data.

Evidence

- Password age: 1,247 days

- Required complexity not met

- Similar credential pattern observed in breach intelligence

- MFA is not enabled for this account

Recommended action

Change the password immediately and enable multi-factor authentication.

Fix difficulty

Low

Sentinel automation

Manual approval required before any account changes are made.

Expert details

Policy: SEC-AUTH-004

Identity object: id:nas-admin-001

Evidence record: ev-identity-auth-2391

Risk score impact: -18 points

8. User Experience Requirements

```
| UX Requirement | Description |
|---|---|
| UX-001 | Every finding shall have a non-technical summary. |
| UX-002 | Every finding shall have an expert-detail expansion. |
| UX-003 | Every recommendation shall identify expected effort. |
| UX-004 | Every recommendation shall identify expected security benefit. |
| UX-005 | Every automated remediation shall require clear user consent unless pre-approved by
policy. |
| UX-006 | Every finding shall be linked to evidence. |
| UX-007 | Every user-facing explanation shall avoid unnecessary fear-based language. |
| UX-008 | The platform shall support home, SMB, enterprise, and expert communication modes. |
| UX-009 | Users shall always be able to understand what data Sentinel used to reach a
conclusion. |
| UX-010 | Users shall always know what to do next. |
```

9. Engineering Impact

These principles are mandatory architectural requirements and apply to:

- All Engineering Archives.

- All Implementation Specifications.
- All coding milestones.
- All user interfaces.
- APIs.
- Reports.
- Dashboards.
- AI-generated explanations.
- Automated recommendations.
- Documentation.
- Future product modules.

Any engine that produces a finding shall implement the Finding Communication Standard.

Any engine that produces remediation guidance shall implement the Evidence Before Opinion principle.

Any user-facing workflow shall implement Simplicity First and Expert Depth on Demand.

10. Acceptance Criteria

This Charter update is accepted when:

- The principles are included in the Project Charter.
- The principles are referenced by future Engineering Archives.
- The Finding Communication Standard is included in implementation guidelines.
- UX requirements are included in the implementation readiness baseline.
- Future coding milestones treat findings and recommendations as structured objects, not plain text.
- Every future UI design supports progressive disclosure from simple explanation to technical evidence.

11. Implementation Guidance for Codex

When implementing Project Sentinel, Codex should treat findings as structured, typed objects.

The implementation should include a canonical `Finding` model with fields for summary, explanation, evidence, recommended actions, severity, confidence, technical details, and audit history.

The UI should never render raw technical findings directly to non-expert users. Instead, raw findings should be transformed into user-appropriate explanations using the communication mode selected for the user or organization.

12. Charter Amendment Status

This document is approved as a Project Sentinel Charter v2.0 amendment and shall be included in the pre-coding baseline.