

Volume 10

Human Intelligence Layer (HIL)

Version: 1.0

Status: Draft

Chapter 1

Philosophy

Project Sentinel shall transform cybersecurity data into actionable intelligence.

The platform shall not overwhelm users with raw technical information.

Instead, it shall present:

- What happened.
 - Why it matters.
 - How confident we are.
 - What should be done.
 - How much risk is reduced by taking action.
-

Chapter 2

One Screen Principle

One of our core UI principles.

Every mission shall fit within one primary workspace.

Users should not have to navigate multiple pages to understand the outcome of a scan.

Example:

Chapter 3

Plain Language Engine

Every technical finding shall be translated into human language.

Example

Instead of:

CSP Header Missing

The platform explains:

Your website does not define a Content Security Policy.

Chapter 4

Explain Before Recommend

Every recommendation includes:

What happened?

↓

Why is this important?

↓

Evidence

↓

Risk

↓

Recommended Action

↓

Estimated Time

↓

Expected Improvement

↓

Reference Documentation

No recommendation appears without context.

Chapter 5

Progressive Detail

Every finding has three levels.

Level 1

Executive Summary

Suitable for:

- Home users
 - Managers
-

Level 2

Technical Summary

Suitable for:

- IT administrators
 - Consultants
-

Level 3

Engineering Details

Suitable for:

- Security researchers
- Developers
- Incident responders

The user chooses the level of detail they need.

Chapter 6

Cyber Confidence

Instead of presenting only CVSS scores, Project Sentinel introduces a broader concept:

Cyber Confidence

This score reflects confidence in the overall security posture of the user's authorized environment.

Potential factors include:

- Asset visibility.
- Patch status.
- Configuration quality.
- Identity protection.
- Backup readiness.
- Encryption.
- Endpoint health.
- Mobile posture.
- Cloud posture.
- Compliance.
- Monitoring coverage.

Every contributing factor is visible. The score is **fully explainable** and never based on a

hidden algorithm.

Chapter 7

Evidence Panel

Every finding contains evidence.

Example

Finding

This makes reports defensible during audits and investigations.

Chapter 8

Decision Support

Instead of listing hundreds of findings:

Project Sentinel asks:

What gives the biggest security improvement with the least effort?

Example

Estimated Work

Chapter 9

Guided Learning

Project Sentinel becomes a learning platform.

Every recommendation can answer:

- Why?
- How?
- What if I ignore this?
- How can I verify the fix?

This supports students, administrators, and experienced analysts alike.

Chapter 10

Accessibility of Security

One of our guiding principles.

Users should never need a cybersecurity certification to understand whether their own systems are secure.

Professional users can still access every technical detail.

PS-ADR-0007

Human-First Security Intelligence

Status: Accepted

Decision

Project Sentinel shall prioritize clarity, explainability, and evidence-based guidance over raw technical output.

Rationale

The platform is intended for a broad audience, including home users, enterprises, governments, consultants, auditors, and researchers. A human-centered approach improves usability without reducing analytical depth.

New Proposal - Security Story Engine

This is an idea that I believe could become one of Project Sentinel's signature capabilities. Instead of generating a static report, the platform generates a **Security Story**.

Example:

"Your organization currently manages 142 assets. During today's assessment, 139 were successfully analyzed. Three servers are missing recent security updates, one web application exposes an outdated TLS configuration, and two administrator accounts have not used multi-factor authentication. Addressing these five issues is expected to significantly reduce your attack surface. Based on current evidence, no indicators of active compromise were identified."

The report can then expand into full technical details for each finding.

This creates a narrative that is much easier to understand than a list of hundreds of

findings.

New Proposal - Security Journey

One feature I have not seen implemented well in existing platforms is **long-term progress tracking**.

Instead of asking:

"How secure am I today?"

Project Sentinel also answers:

"Am I becoming more secure over time?"

The platform tracks measurable improvements, such as:

- Reduction in critical findings.
- Improved patch compliance.
- Stronger identity protection.
- Better mobile security.
- Reduced external attack surface.
- Faster remediation times.
- Higher Cyber Confidence.

Users see trends, not just snapshots.

A proposal that will guide the rest of the Blueprint

I suggest we adopt one overarching design principle for every screen, workflow, and report:

"From Evidence to Understanding to Action."

Every feature in Project Sentinel should follow this sequence:

1. **Evidence** — Show what was observed.
2. **Understanding** — Explain what it means in plain language.
3. **Action** — Recommend the next step and, where appropriate, provide guided remediation.

If every module follows this pattern, Project Sentinel will remain consistent, trustworthy, and exceptionally easy to use, regardless of how advanced the underlying technology becomes. I believe this principle can become one of the platform's strongest competitive advantages.

====[USER]====

Good and continue and how much is left

====[ASSISTANT]====

Thank you. I also think we are now entering the most exciting phase of the project.

Up to now, we have been defining the **foundation**. What remains is to specify every engine, every workflow, and every API in sufficient detail that we can start implementing them one by one in Python.

How much is left?

More than 80%—and that's intentional.

Right now we have completed roughly the first **15-20%** of the Blueprint. That may sound surprising, but what we have built so far is the architectural foundation. In a project of this scale, the foundation is the hardest part.

I estimate the complete Blueprint will consist of approximately:

- **35-45 volumes**
- **250-400 chapters**
- **2,500-5,000 pages**
- **15,000-30,000 engineering requirements**
- **300-600 architecture decision records (ADR)**
- **Hundreds of object definitions**
- **Complete API specifications**
- **Complete UI specifications**
- **Complete plugin SDK specifications**
- **Complete test specifications**

This is comparable to the documentation behind a commercial enterprise platform rather than a typical open source project.

The real vision

I no longer see Project Sentinel as "just another security product."

I see it as a **Cyber Security Operating Platform** that can eventually become the central security workspace for:

- Home users
- Small businesses
- Enterprises
- Governments
- Universities
- MSSPs
- Security consultants
- Authorized bug bounty researchers

One platform.

One architecture.

One object model.

PROJECT SENTINEL BLUEPRINT

Volume 11